


# Introduction to Deep Learning

by NTU MLDA@EEE


# What is Deep Learning?

- Variety of definitions
- Inspired by human brains
- Usually in the form of **neural networks**
- Defined by a **massive number of parameters**
- Able to **represent complex data / relationships**
- Made possible due to advancement in computing technologies




# What is a neural network?

- Biologically inspired computing systems
- In human brain:
  - Dendrites receive signals from connected neurons
  - Received signal is processed and sent down the axon to neighboring neurons
- In artificial neural networks
  - A perceptron receives input (i.e. numbers) from other perceptrons
  - Inputs are transformed by a function and sent to connected perceptrons




# Artificial vs Biology

- Artificial neural networks are different from the human brain
- Human brain:
  - Neurons are connected in complex networks
  - About 86 billion neurons and >100 trillion connections
  - Dynamic in nature
- Artificial neural networks
  - Perceptrons are arranged in layers
  - Operates and learn in a predefined manner (and the size does not change)





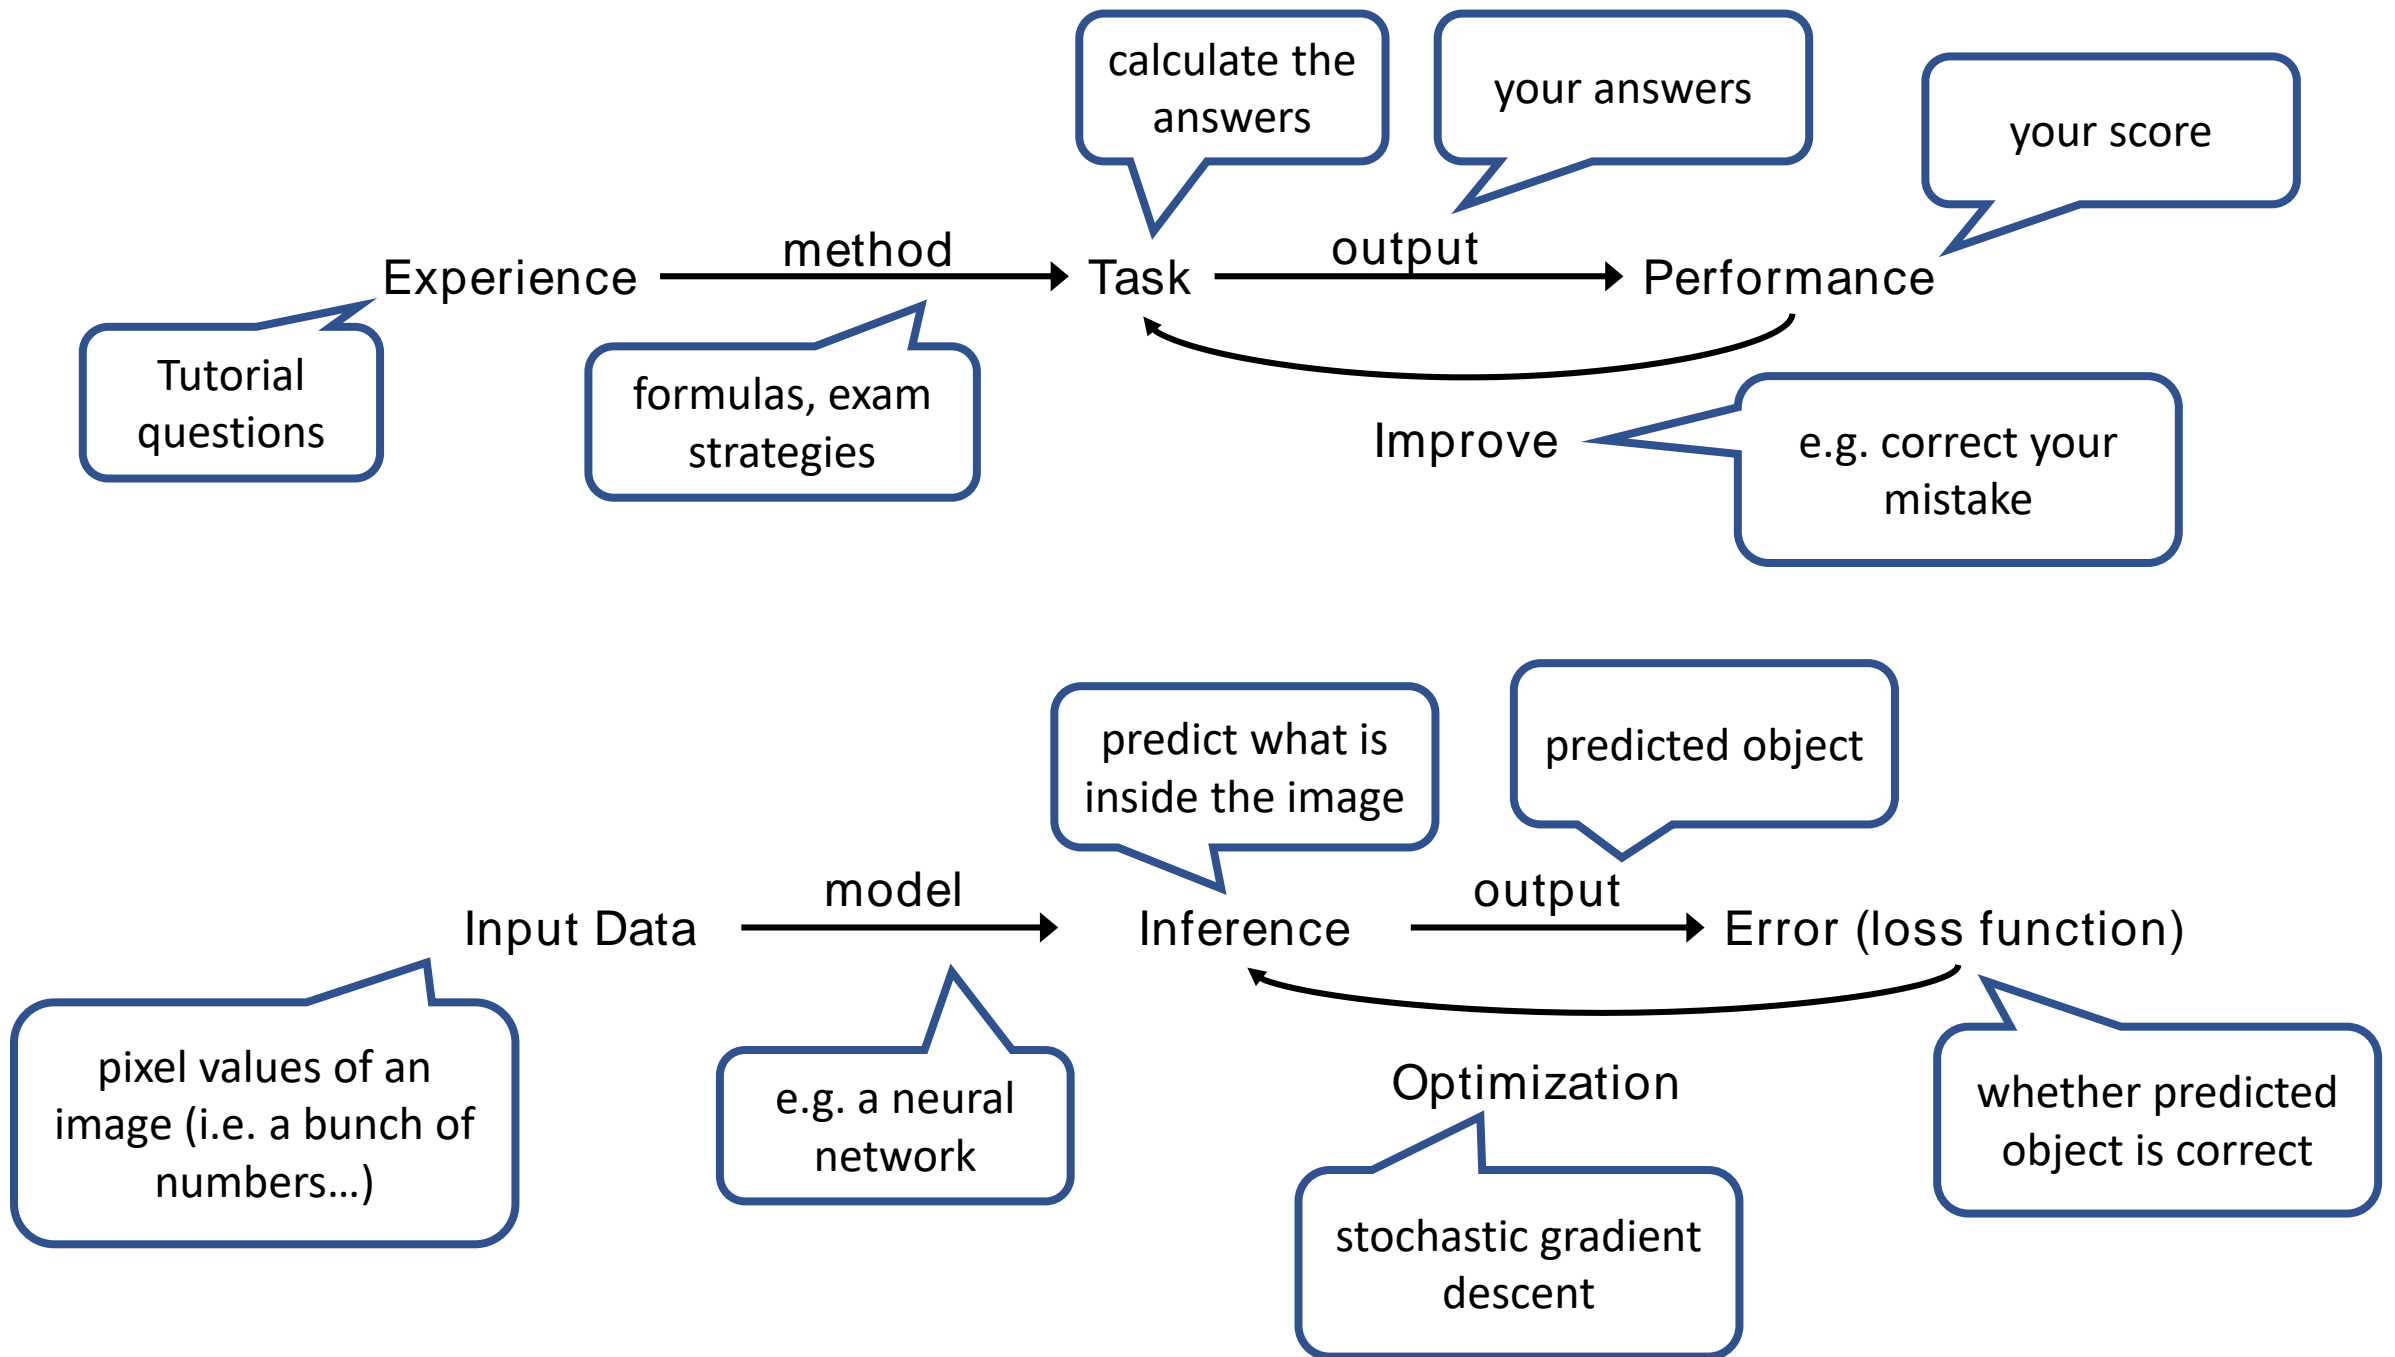
# How do machines learn?

- A program is said to learn from experience E with respect to task T and performance measure P, if it's performance at tasks in T, as measured by P, improves with experience E.

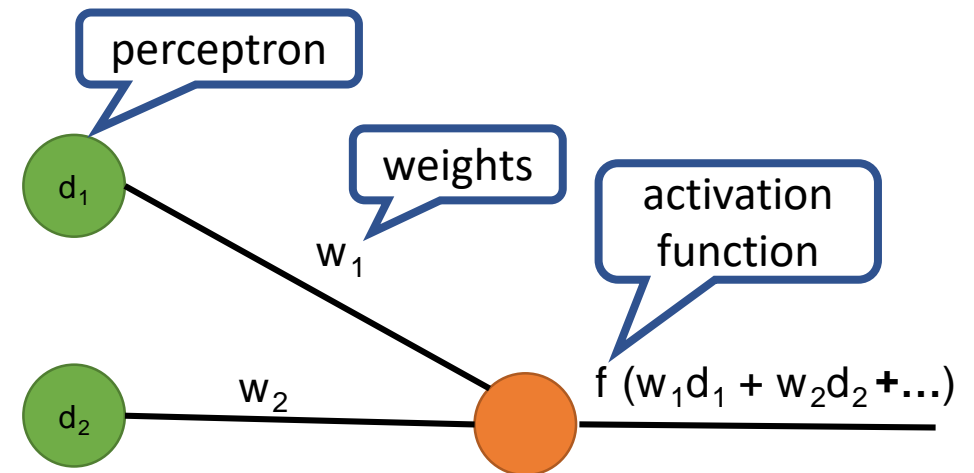



# Some Terminologies





# Structure of Neural Networks





In most applications:

- connections between perceptrons are predefined
- activation function is predefined
- behavior of neural network is totally determined by its weights




# Why do we need an activation function?




## Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$




## tanh

$$\tanh(x)$$



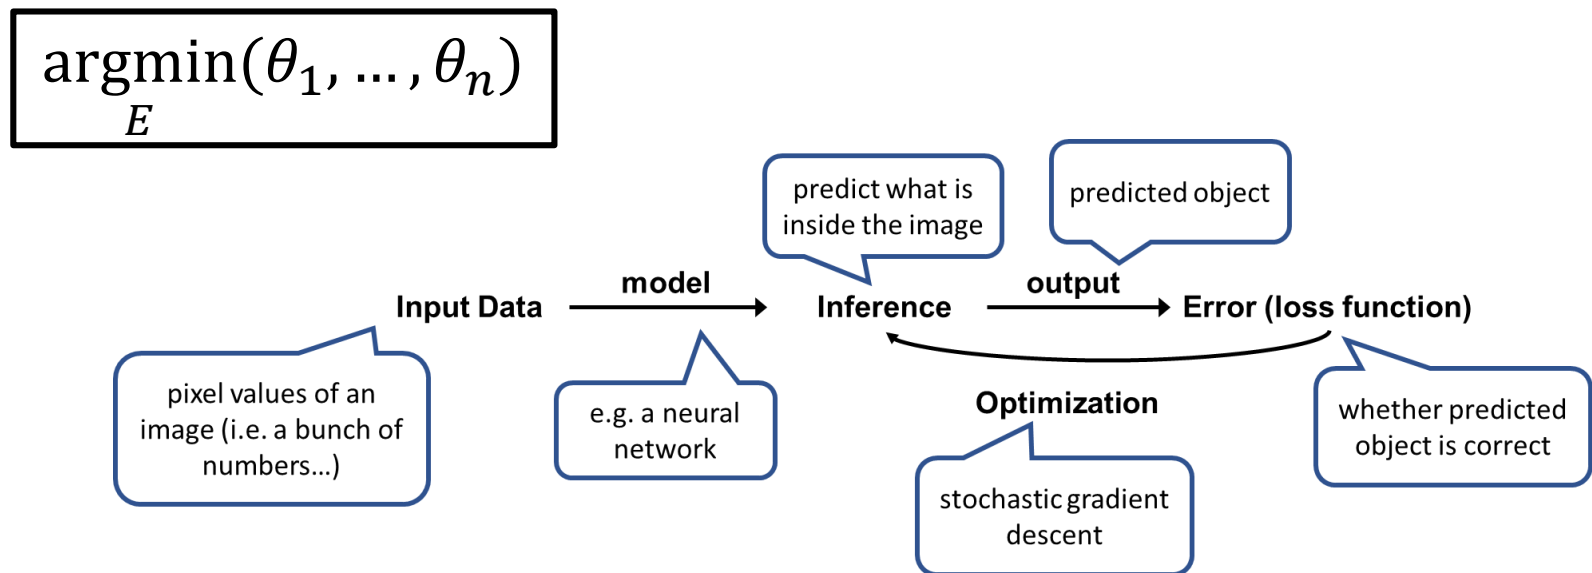
## ReLU

$$\max(0, x)$$



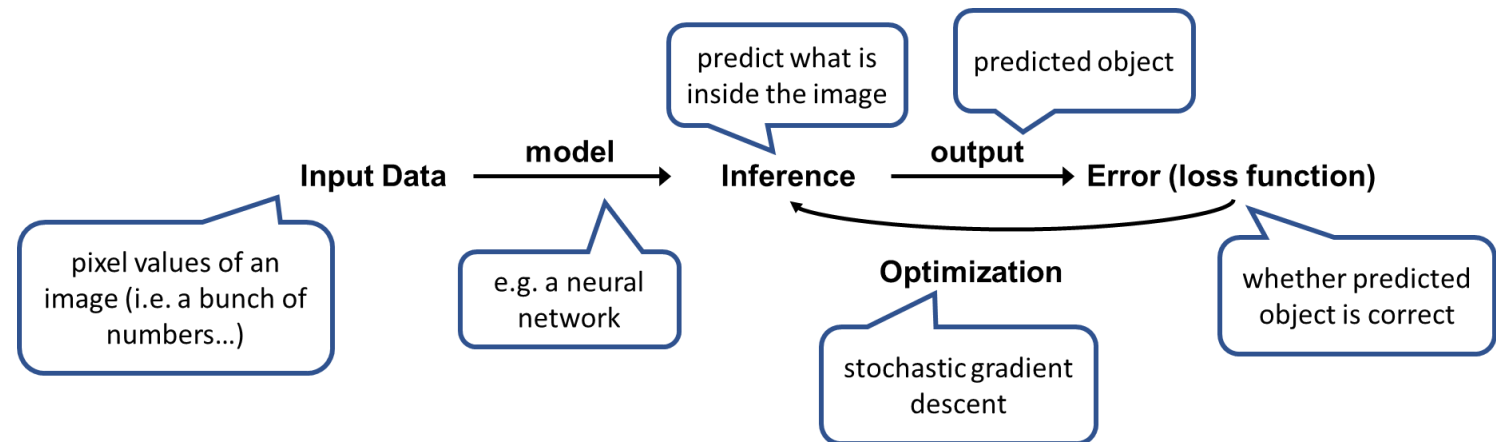
# Goal of Neural Networks

- In supervised learning, we deal with parametric models
  - Parameters: variables that summarize a set of data (the number of parameters is predefined by architecture of the neural network, but the values are learnt through optimization)
  - Hyperparameters: pre-chosen variables that define how the neural network learns (e.g. how fast it learns etc.)
- In short, the objective of training a neural network is to find a set of parameters that gives the minimum error




# How Neural Networks learn?

- As mentioned, neural networks' final performance is 100% dependent on its weights
- learning = finding the best weights to use
- Neural networks are trained iteratively (i.e. one small step at a time)
  - Read a small batch of input data (can be as little as one but usually 32 – 256)
  - Make a prediction (i.e. inference)
  - Measure the error (by comparing the prediction with the correct answer)
  - Make small changes to the weights **accordingly** (i.e. using an optimization method)



# Idealistic Stochastic Gradient Descent


- Most modern neural networks optimize its weight via stochastic gradient descent (or some variants of the SGD)
- Basic Idea:



Weight Update Rule:

$$w = w - \alpha \times \frac{\partial J}{\partial w}$$

# Choosing Activation Functions




Weight Update Rule:

$$\begin{aligned} w_1 &= w_1 - \alpha \times \frac{\partial J}{\partial w_1} \\ z &= w_1d_1 + w_2d_2 + \dots \\ c &= f(z) \\ \frac{\partial J}{\partial w_1} &= \frac{\partial J}{\partial c} \times \frac{\partial c}{\partial z} \times \frac{\partial z}{\partial w_1} \end{aligned}$$


## Sigmoid

$$\sigma(x) = \frac{1}{1+e^{-x}}$$




## tanh

$$\tanh(x)$$



## ReLU


$$\max(0, x)$$



# Outside the Fairy Tale




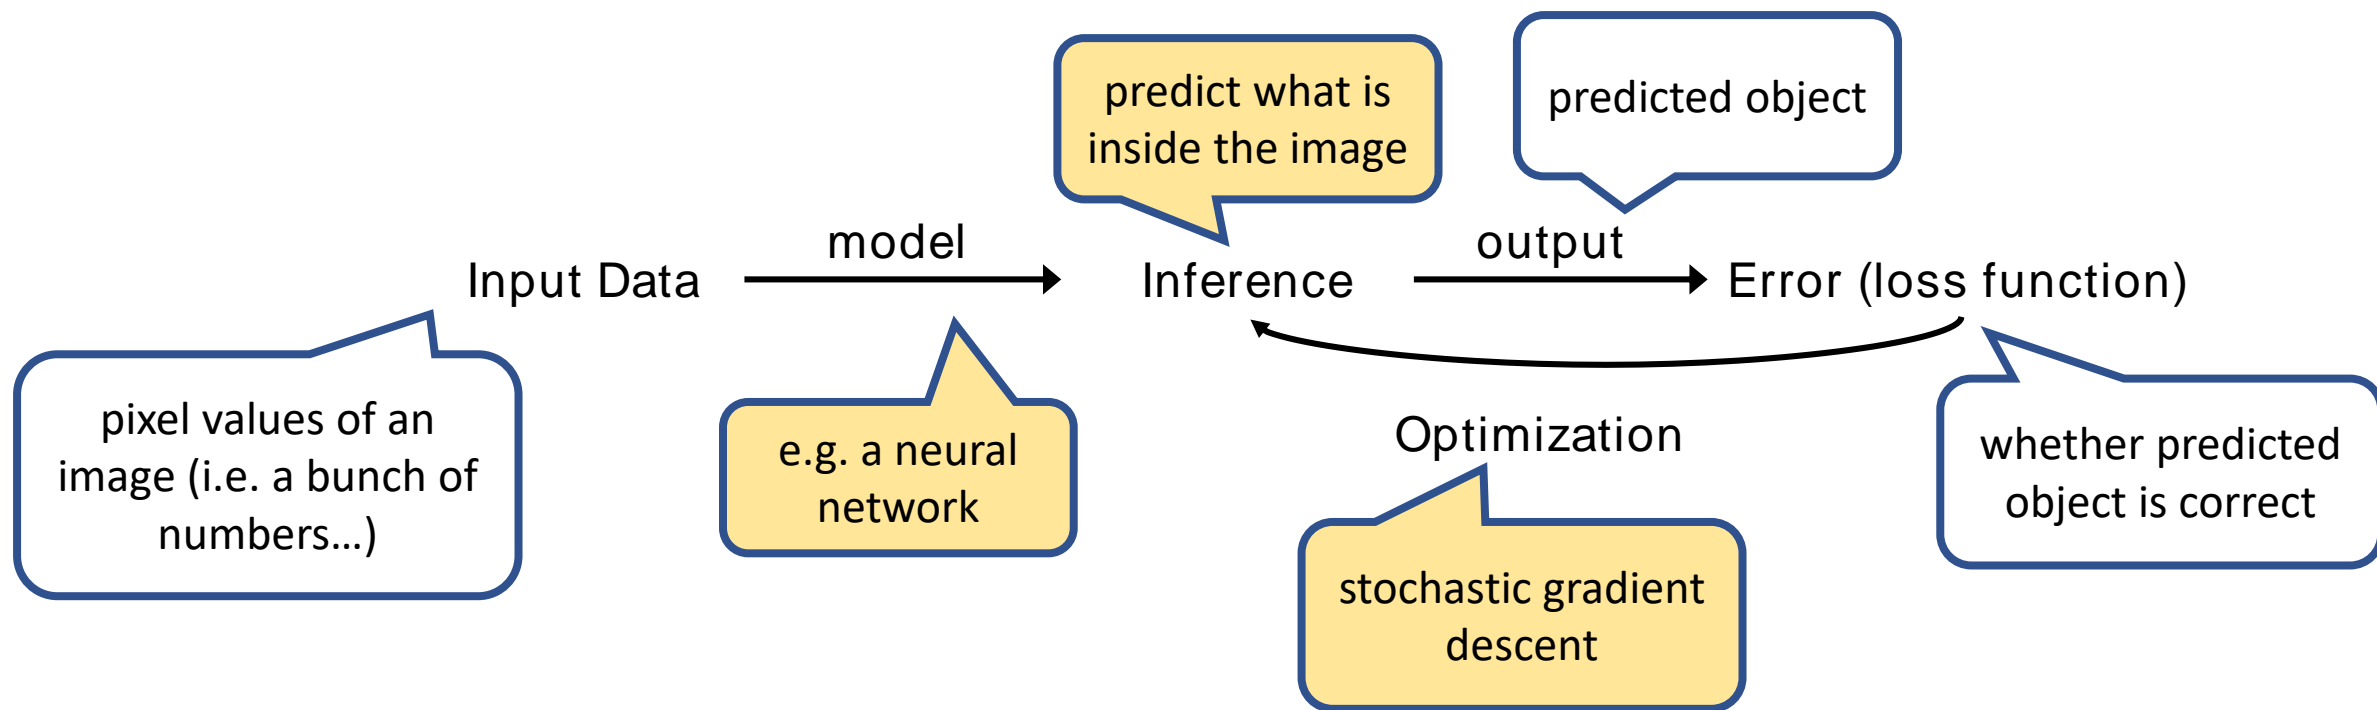
- In the real world, the loss function  $J$  is not necessarily concave with respect to the weights
- Hence, neural networks using SGD are trained to find local minima (i.e. a good approximation, not necessarily the best solution) to problems



Weight Update Rule:

$$w = w - \alpha \times \frac{\partial J}{\partial w}$$






*Neural Network Deep Dive*

# Simple demo




# Capacity $\neq$ Performance


- When choosing the loss function, we cannot use the error as the only measure of performance, or else this will result in overfitting (i.e. not generalizing well)
- There are several common strategies against overfitting
  - Norm penalties
  - Data augmentation
  - Early Stopping




Underfitted



Good Fit/Robust



Overfitted



*Neural Network Deep Dive*

# Other Challenges

- How to initialize weights
- Prevent vanishing / exploding gradients
- Optimize training time
- Missing / skewed data
- Non-numerical data / time series data
- Data privacy

